

Evidence theory on prediction error

Zhuoyan Xu

Aug 10 2019

Consider we have an extra covariate matrix $X^{d_1 \times p}$ (accounting for features), which contains the information of countries. We want to connect the membership matrix (or factor matrix) A and B with the information in tensor X.

The general form is:

$$\begin{aligned} \text{logit} \{ \mathbb{E} [\mathcal{Y}^{d_1 d_2 \dots d_K}] \} &= \Theta = \mathcal{G}^{r_1 r_2 \dots r_K} \times_1 W_1^{d_1 r_1} \times_2 W_2^{d_2 r_2} \times_3 N_3^{d_2 r_2} \dots \times_K N_K^{d_K r_K} \\ W_1^{d_1 r_1} &= X_1^{d_1 p} N_1^{p r_1} \\ W_2^{d_2 r_2} &= X_1^{d_2 p} N_2^{p r_2} \end{aligned}$$

where \mathcal{G} is the low rank core tensor of factorization. W_1, W_2, \dots, N_K are factor matrices. Without out loss of generality, N_i is the regression coefficient matrix for X_i on W_i .

We can write down the model in another view, which helps to compute:

$$\begin{aligned} \Theta &= \mathcal{C} \times_1 X_1 \times_2 X_2 \\ \mathcal{C} &= \mathcal{G}^{r_1 r_2 \dots r_K} \times_1 N_1^{d_1 r_1} \times_2 N_2^{d_2 r_2} \dots \times_K N_K^{d_K r_K} \end{aligned}$$

where \mathcal{C} , a tensor with tucker rank (r_1, \dots, r_K) , is our coefficient tensor. X_i is our predictor.

1 Frobenius Norm Loss

Without loss of generality, we choose logit as link function, use $\pi = f(\Theta) = \text{logit}^{-1}(\Theta)$ to denote true probability of our tensor.

The log-likelihood function is:

$$\mathcal{L}_Y(\pi) = \sum_{i_1, \dots, i_K} \left[\mathbf{1}_{\{y_{i_1, \dots, i_K}=1\}} \log(\pi_{i_1, \dots, i_K}) + \mathbf{1}_{\{y_{i_1, \dots, i_K}=0\}} \log\{\pi_{i_1, \dots, i_K}\} \right]$$

Thus we have:

$$\begin{aligned} \frac{\partial \mathcal{L}_Y}{\partial \pi_{i_1, \dots, i_K}} &= \frac{1}{\pi_{i_1, \dots, i_K}} \mathbf{1}_{\{y_{i_1, \dots, i_K}=1\}} - \frac{1}{1 - \pi_{i_1, \dots, i_K}} \mathbf{1}_{\{y_{i_1, \dots, i_K}=0\}} \\ \frac{\partial \mathcal{L}_Y^2}{\partial \pi_{i_1, \dots, i_K}^2} &= - \frac{\mathbf{1}_{\{y_{i_1, \dots, i_K}=1\}}}{\pi_{i_1, \dots, i_K}^2} - \frac{\mathbf{1}_{\{y_{i_1, \dots, i_K}=0\}}}{(1 - \pi_{i_1, \dots, i_K})^2} \\ \frac{\partial \mathcal{L}_Y^2}{\partial \pi_{i_1, \dots, i_K} \pi_{i'_1, \dots, i'_K}} &= 0 \quad \text{if} \quad (i_1, \dots, i_K) \neq (i'_1, \dots, i'_K) \end{aligned}$$

Without loss of generality, for any observation y in binary tensor and its corresponding probability π , we have:

$$\frac{\partial \mathcal{L}_y}{\partial \pi} = \frac{y}{\pi} - \frac{1-y}{1-\pi} \leq \frac{1}{\pi(1-\pi)}$$

Since the constrain $\theta = \text{logit}(\pi) \leq \alpha$ (constrain on the max norm of ground truth tensor) and the symmetric property of link function f , we have:

$$\pi(1-\pi) = f(\theta)(1-f(\theta)) \geq f(\alpha)(1-f(\alpha)) = c$$

where c is a constant. Thus:

$$\frac{\partial \mathcal{L}_y}{\partial \pi} \leq \frac{1}{c}$$

similarly, consider

$$\frac{\partial \mathcal{L}_y^2}{\partial \pi^2} = -\frac{y}{\pi^2} - \frac{1-y}{(1-\pi)^2} \leq -1$$

Define

$$\mathcal{S}_y(\pi_{\text{true}}) = \left[\left[\frac{\partial \mathcal{L}_y}{\partial \pi_{i_1, \dots, i_K}} \right] \right]_{\pi=\pi_{\text{true}}} \quad \text{and} \quad \mathcal{H}_y(\pi_{\text{true}}) = \left\| \frac{\partial \mathcal{L}_y^2}{\partial \pi_{i_1, \dots, i_K} \partial \pi_{i'_1, \dots, i'_K}} \right\|_{\pi=\pi_{\text{true}}}$$

Use Taylor expansion and we have:

$$\mathcal{L}_y(\pi) = \mathcal{L}_y(\pi_{\text{true}}) + \langle \mathcal{S}_y(\pi_{\text{true}}), \pi - \pi_{\text{true}} \rangle + \frac{1}{2} \text{vec}(\pi - \pi_{\text{true}})^T \mathcal{H}_y(\tilde{\pi}) \text{vec}(\pi - \pi_{\text{true}})$$

Use $\hat{\pi}$ to denote MLE prediction, we have:

$$\begin{aligned} 0 \leq \mathcal{L}_y(\pi) - \mathcal{L}_y(\pi_{\text{true}}) &= \langle \mathcal{S}_y(\pi_{\text{true}}), \pi - \pi_{\text{true}} \rangle + \frac{1}{2} \text{vec}(\pi - \pi_{\text{true}})^T \mathcal{H}_y(\tilde{\pi}) \text{vec}(\pi - \pi_{\text{true}}) \\ &\leq \langle \mathcal{S}_y(\pi_{\text{true}}), \pi - \pi_{\text{true}} \rangle - \frac{1}{2} \|\pi - \pi_{\text{true}}\|_F^2 \end{aligned}$$

Apply the theorem on Gaussian width, we have our prediction error:

$$\|\hat{\pi} - \pi\|_F \leq \frac{2C_2}{f(\alpha)(1-f(\alpha))} \sqrt{\prod_{k=1}^{K-1} r_k \sum_{k=1}^K d_k}$$

2 KL-Divergence

Consider $\pi = f(\Theta)$, where f is the link function. We use Without loss of generality, for any observation y_i in binary tensor and its corresponding ground truth θ_i and probability π_i , we have:

$$\mathcal{L}_y(\Theta) = \sum_i \{y_i \log[f(\theta)] + (1-y_i) \log[1-f(\theta)]\}$$

Use Θ^* to denote the true parameter, then we have:

$$\mathcal{L}_{\mathcal{Y}}(\Theta) - \mathcal{L}_{\mathcal{Y}}(\Theta^*) = \sum_i \left\{ y_i \log \left[\frac{f(\theta)}{f(\theta^*)} \right] + (1 - y_i) \log \left[\frac{1 - f(\theta)}{1 - f(\theta^*)} \right] \right\}$$

Thus, we have:

$$\begin{aligned} \mathbb{E}_{\Theta^*} \{ \mathcal{L}_{\mathcal{Y}}(\Theta) - \mathcal{L}_{\mathcal{Y}}(\Theta^*) \} &= \sum_{y_i=1} \left\{ f(\theta^*) \log \left[\frac{f(\theta)}{f(\theta^*)} \right] \right\} + \sum_{y_i=0} \left\{ [1 - f(\theta^*)] \log \left[\frac{1 - f(\theta)}{1 - f(\theta^*)} \right] \right\} \\ &= \sum_i \left\{ \pi^* \log \left[\frac{\pi}{\pi^*} \right] \right\} \\ &= -D_{\text{KL}}(\pi^* \parallel \pi) \end{aligned}$$

Thus we have:

$$\mathbb{E}_{\pi^*} \{ \mathcal{L}_{\mathcal{Y}}(\pi) - \mathcal{L}_{\mathcal{Y}}(\pi^*) \} = -D_{\text{KL}}(\pi^* \parallel \pi)$$

Recalling Taylor expansion in previous section:

$$\mathcal{L}_{\mathcal{Y}}(\pi) = \mathcal{L}_{\mathcal{Y}}(\pi^*) + \langle S_{\mathcal{Y}}(\pi^*), \pi - \pi^* \rangle + \frac{1}{2} \text{vec}(\pi - \pi^*)^T \mathcal{H}_{\mathcal{Y}}(\check{\pi}) \text{vec}(\pi - \pi^*)$$

Recalling the definition of $S_{\mathcal{Y}}(\pi^*)$ and $\mathcal{H}_{\mathcal{Y}}(\check{\pi})$, we have

$$\mathbb{E}_{\pi^*} S_{\mathcal{Y}}(\pi^*) = 0$$

And for $\forall \pi_i$, we have:

$$-\mathbb{E}_{\pi^*} \left[\frac{\partial \mathcal{L}_{\mathcal{Y}}^2}{\partial \pi_i^2} \right] = \mathbb{E}_{\pi^*} \left[\frac{y}{\pi_i^2} + \frac{1 - y}{(1 - \pi_i)^2} \right] = \left[\frac{\pi^*}{\pi_i^2} + \frac{1 - \pi^*}{(1 - \pi_i)^2} \right]$$

Similar to max norm on ground truth tensor Θ , we have

$$-\mathbb{E}_{\pi^*} \left[\frac{\partial \mathcal{L}_{\mathcal{Y}}^2}{\partial \pi_i^2} \right] = \left[\frac{\pi^*}{\pi_i^2} + \frac{1 - \pi^*}{(1 - \pi_i)^2} \right] \leq \frac{2}{\pi_i^2(1 - \pi_i)^2} \leq \frac{2}{f(\alpha)^2(1 - f(\alpha))^2}$$

Thus:

$$\begin{aligned} D_{\text{KL}}(\pi^* \parallel \pi) &= -\mathbb{E}_{\pi^*} \{ \mathcal{L}_{\mathcal{Y}}(\pi) - \mathcal{L}_{\mathcal{Y}}(\pi^*) \} \\ &= -\frac{1}{2} \mathbb{E}_{\pi^*} \{ \text{vec}(\pi - \pi^*)^T \mathcal{H}_{\mathcal{Y}}(\check{\pi}) \text{vec}(\pi - \pi^*) \} \\ &\leq \frac{1}{f(\alpha)^2(1 - f(\alpha))^2} \|\hat{\pi} - \pi\|_F^2 \end{aligned}$$

According to our result in previous section, we have:

$$D_{\text{KL}}(\pi^* \parallel \pi) \leq \frac{1}{f(\alpha)^2(1 - f(\alpha))^2} \|\hat{\pi} - \pi\|_F^2 \leq \frac{4C_2^2}{f(\alpha)^4(1 - f(\alpha))^4} \prod_{k=1}^{K-1} r_k \sum_{k=1}^K d_k$$

3 Hellinger Loss

According to the definition of Hellinger loss, we have:

$$\begin{aligned}
d_H^2(\pi, \pi^*) &= \sum_{i=1} \left\{ (\sqrt{\pi_i} - \sqrt{\pi_i^*})^2 + (\sqrt{1 - \pi_i} - \sqrt{1 - \pi_i^*})^2 \right\} \\
&\leq \sum_{i=1} \left\{ |\sqrt{\pi_i} - \sqrt{\pi_i^*}|(\sqrt{\pi_i} + \sqrt{\pi_i^*}) + |\sqrt{1 - \pi_i} - \sqrt{1 - \pi_i^*}|(\sqrt{1 - \pi_i} + \sqrt{1 - \pi_i^*}) \right\} \\
&= 2 \sum_{i=1} |\pi_i - \pi_i^*|
\end{aligned}$$

According to Pinsker's inequality, we have:

$$\sum_{i=1} |\pi_i - \pi_i^*| \leq \sqrt{\frac{1}{2} D_{\text{KL}}(P \| Q)}$$

Thus, we have:

$$d_H^2(\pi, \pi^*) \leq \sqrt{2 D_{\text{KL}}(P \| Q)} \leq \frac{2C_2}{f(\alpha)^2(1 - f(\alpha))^2} \sqrt{2 \prod_{k=1}^{K-1} r_k \sum_{k=1}^K d_k}$$